

1. **What is the primary goal of horizontal scaling in a distributed system?**

- a. Increase the capacity of a single machine
- b. Distribute the load across multiple machines
- c. Upgrade the hardware components of a machine
- d. Optimize the software architecture

Answer: b. Distribute the load across multiple machines

2. **Which scaling approach is more likely to involve the use of commodity hardware?**

- a. Vertical scaling
- b. Horizontal scaling
- c. Both use similar types of hardware
- d. Neither uses commodity hardware

Answer: b. Horizontal scaling

3. **In vertical scaling, what component is typically upgraded to handle increased demand?**

- a. Storage capacity
- b. Network bandwidth
- c. Number of machines
- d. Resources like CPU, RAM, or storage

Answer: d. Resources like CPU, RAM, or storage

4. **What is a key advantage of horizontal scaling in terms of fault tolerance?**

- a. It is less expensive
- b. It doesn't require changes to the system architecture
- c. The failure of one machine doesn't bring down the entire system

d. It is easier to implement initially

Answer: c. The failure of one machine doesn't bring down the entire system

5. Which scaling strategy is more likely to involve adding servers to a server farm to handle increased traffic?

a. Vertical scaling

b. Horizontal scaling

c. Both are equally likely

d. Neither involves adding servers

Answer: b. Horizontal scaling

6. What is a potential drawback of vertical scaling?

a. Difficulty in initial implementation

b. Limited by the capacity of a single machine

c. Requires changes to the software architecture

d. Higher overall cost

Answer: b. Limited by the capacity of a single machine

7. Which scaling approach is more suitable for applications with unpredictable and varying workloads?

a. Vertical scaling

b. Horizontal scaling

c. Both are equally suitable

d. Neither is suitable

Answer: b. Horizontal scaling

8. In which scenario is vertical scaling often preferred?

- a. When cost is a critical factor
- b. When fault tolerance is a top priority
- c. When the system needs to handle a rapidly increasing number of users
- d. When the application is designed for a distributed environment

Answer: b. When fault tolerance is a top priority

9. What term is commonly used to describe the ability of a system to handle increased loads by adding more machines?

- a. Load balancing
- b. Capacity planning
- c. Scalability
- d. Elasticity

Answer: c. Scalability

10. Which scaling approach is more likely to require a careful analysis of the system architecture and potential redesigns?

- a. Vertical scaling
- b. Horizontal scaling
- c. Both require similar analysis
- d. Neither requires architectural analysis

Answer: a. Vertical scaling

1. What is the primary goal of database normalization?

- a. Maximizing redundancy
- b. Minimizing data consistency

- c. Minimizing data redundancy
- d. Maximizing data consistency

Answer: c. Minimizing data redundancy

2. **In which normal form are all non-prime attributes functionally dependent on the primary key, the whole primary key, and nothing but the primary key?**
- a. First Normal Form (1NF)
 - b. Second Normal Form (2NF)
 - c. Third Normal Form (3NF)
 - d. Boyce-Codd Normal Form (BCNF)

Answer: a. First Normal Form (1NF)

3. **Which normal form deals with the removal of transitive dependencies?**
- a. First Normal Form (1NF)
 - b. Second Normal Form (2NF)
 - c. Third Normal Form (3NF)
 - d. Fourth Normal Form (4NF)

Answer: c. Third Normal Form (3NF)

4. **What is a candidate key in a relation?**
- a. Any attribute that is part of the primary key
 - b. A minimal superkey for uniquely identifying tuples
 - c. A non-key attribute in a relation
 - d. A foreign key in another table

Answer: b. A minimal superkey for uniquely identifying tuples

5. **In Second Normal Form (2NF), what condition must be satisfied in addition to First Normal Form (1NF)?**
- a. No partial dependencies
 - b. No transitive dependencies

- c. All attributes are atomic
- d. All attributes are integers

Answer: a. No partial dependencies

6. **Which normal form ensures that there are no non-prime attribute dependencies on superkeys?**
- a. Boyce-Codd Normal Form (BCNF)
 - b. Fourth Normal Form (4NF)
 - c. Fifth Normal Form (5NF)
 - d. Sixth Normal Form (6NF)

Answer: a. Boyce-Codd Normal Form (BCNF)

7. **In normalization, what is the process of decomposing a relation into smaller, well-structured relations without loss of information called?**
- a. Normalization
 - b. Denormalization
 - c. Decomposition
 - d. Aggregation

Answer: c. Decomposition

8. **Which normal form deals with multivalued dependencies?**
- a. Third Normal Form (3NF)
 - b. Fourth Normal Form (4NF)
 - c. Fifth Normal Form (5NF)
 - d. Sixth Normal Form (6NF)

Answer: b. Fourth Normal Form (4NF)

9. **What is the purpose of denormalization in a database design?**
- a. Minimizing data redundancy
 - b. Improving data consistency

- c. Maximizing data redundancy for faster retrieval
- d. Reducing the number of tables in a database

Answer: c. Maximizing data redundancy for faster retrieval

10. Which normal form ensures that there are no join dependencies?

- a. Fourth Normal Form (4NF)
- b. Fifth Normal Form (5NF)
- c. Sixth Normal Form (6NF)
- d. Domain-Key Normal Form (DKNF)

Answer: d. Domain-Key Normal Form (DKNF)

Indexing Tabular Data:

1. What is indexing in the context of tabular data?

- a. A way to organize data in columns
- b. A technique to create a new table
- c. A mechanism to optimize data retrieval based on specific columns
- d. A process of normalizing data

Answer: c. A mechanism to optimize data retrieval based on specific columns

2. Which of the following is a key advantage of indexing tabular data?

- a. Reduced storage space
- b. Faster data retrieval for specific columns
- c. Improved data consistency
- d. Enhanced data encryption

Answer: b. Faster data retrieval for specific columns

3. In a tabular data structure, what is an index typically associated with?

- a. Entire rows of data
- b. Specific columns of data
- c. Primary key constraints
- d. Foreign key constraints

Answer: b. Specific columns of data

4. What type of query benefits the most from indexing tabular data?

- a. Queries involving multiple joins
- b. Queries filtering or sorting based on specific columns
- c. Queries involving aggregate functions
- d. Queries with subqueries

Answer: b. Queries filtering or sorting based on specific columns

Indexing Columnar Data:

5. What is a characteristic of columnar databases?

- a. Organizing data in rows
- b. Optimizing for fast transaction processing
- c. Storing data in columnar format rather than row-wise
- d. Using primary keys for data retrieval

Answer: c. Storing data in columnar format rather than row-wise

6. What advantage does indexing in a columnar database provide over traditional row-based databases?

- a. Improved transaction processing speed
- b. Reduced storage space
- c. Faster data retrieval for specific columns

d. Enhanced data consistency

Answer: c. Faster data retrieval for specific columns

7. Which type of data is well-suited for columnar indexing?

- a. Data with high update frequency
- b. Data with a high degree of data redundancy
- c. Data with analytical or reporting queries
- d. Data with complex relationships

Answer: c. Data with analytical or reporting queries

8. In a columnar database, what is typically indexed for faster query performance?

- a. Entire rows of data
- b. Specific columns of data
- c. Primary key constraints
- d. Foreign key constraints

Answer: b. Specific columns of data

General Indexing:

9. What is a common challenge associated with indexing regardless of the data structure?

- a. Increased storage space
- b. Slower data retrieval
- c. Overhead in terms of insert, update, and delete operations
- d. Limited support for query optimization

Answer: c. Overhead in terms of insert, update, and delete operations

10. In which scenario might a combination of row-based and columnar indexing be beneficial?

- a. OLTP (Online Transaction Processing) systems
- b. OLAP (Online Analytical Processing) systems
- c. Systems with low query complexity

d. Systems with frequent data updates

Answer: b. OLAP (Online Analytical Processing) systems

Linux File System MCQ

1. Which file system is commonly used for the root file system in most Linux distributions?

a. ext2

b. ext3

c. ext4

d. XFS

Answer: c. ext4

2. What is the purpose of the `/bin` directory in Linux?

a. Temporary files

b. System binaries

c. User binaries

d. Configuration files

Answer: b. System binaries

3. Which command is used to display the current working directory in Linux?

a. `pwd`

b. `cd`

c. `dir`

d. `ls`

Answer: a. `pwd`

4. In Linux, what does the command `chmod 755 file.txt` do?

a. Adds read and write permissions for the owner

- b. Adds execute permissions for the group
- c. Adds read, write, and execute permissions for the owner
- d. Adds read and execute permissions for everyone

Answer: d. Adds read and execute permissions for everyone

5. Which file system type is often used for removable media like USB drives in Linux?

- a. FAT32
- b. ext4
- c. NTFS
- d. HFS+

Answer: a. FAT32

6. What is the purpose of the `/dev` directory in Linux?

- a. Device files
- b. System configuration files
- c. User home directories
- d. Temporary files

Answer: a. Device files

7. Which command is used to create an empty file in Linux?

- a. `create`
- b. `touch`
- c. `newfile`
- d. `make`

Answer: b. `touch`

8. In Linux, what is the function of the `mount` command?

- a. Create a new directory
- b. Attach a file system to the directory tree

- c. Display disk usage statistics
- d. Change file permissions

Answer: b. Attach a file system to the directory tree

9. Which directory in Linux contains system log files?

- a. /var/log
- b. /log
- c. /var/system
- d. /system/log

Answer: a. /var/log

10. What does the `df -h` command display in Linux?

- a. Directory files
- b. Disk space usage in a human-readable format
- c. Device files
- d. Directory hierarchy

Answer: b. Disk space usage in a human-readable format

11. Which file in Linux contains the configuration information for the GRUB bootloader?

- a. /boot/grub.cfg
- b. /etc/bootloader.conf
- c. /etc/grub.conf
- d. /boot/grub/grub.cfg

Answer: a. /boot/grub.cfg

12. What is the purpose of the `umask` command in Linux?

- a. Display user information
- b. Set file permissions for new files
- c. Change the user's password

d. Mount a file system

Answer: b. Set file permissions for new files

13. Which command is used to remove a directory in Linux?

a. `rmdir`

b. `delete`

c. `rm`

d. `remove`

Answer: a. `rmdir`

14. In Linux, what is the purpose of the `grep` command?

a. Search for files and directories

b. Search for text patterns in files

c. Display the contents of a file

d. Copy files and directories

Answer: b. Search for text patterns in files

15. What is the purpose of the `/etc/passwd` file in Linux?

a. Stores user passwords

b. Configures system-wide settings

c. Lists user accounts and their attributes

d. Manages network configurations

Answer: c. Lists user accounts and their attributes

16. Which command is used to create a symbolic link in Linux?

a. `symlink`

b. `ln -s`

c. `link`

d. `create-link`

Answer: b. ln -s

17. In Linux, what is the purpose of the /tmp directory?

- a. Temporary files
- b. System binaries
- c. Configuration files
- d. User home directories

Answer: a. Temporary files

18. Which file in Linux contains information about system startup processes?

- a. /etc/boot.conf
- b. /boot/startup.log
- c. /var/init.log
- d. /etc/inittab

Answer: d. /etc/inittab

19. What does the echo command do in Linux?

- a. Display disk usage statistics
- b. Print the contents of a file
- c. Display the current date and time
- d. Print a message to the terminal

Answer: d. Print a message to the terminal

20. Which command is used to find files in Linux based on various criteria such as name or size?

- a. search
- b. locate
- c. find
- d. grep

Answer: c. find